

Elementary Statistics: Math 080

Jordan Hanson

August 3, 2022

Whittier College Department of Physics and Astronomy

Summary

Unit 2 and 3

1. Central Limit Theorem: 7.1
2. Confidence Intervals and Hypothesis Testing
 - Confidence intervals and data interpretation: 8.1 - 8.4
 - Rejecting the null hypothesis, types of error, underlying distributions: 9.1 - 9.3

The Central Limit Theorem

The Central Limit Theorem

Central Limit Theorem: Let X be a continuous random variable, with mean μ_X and standard deviation σ_X . The average \bar{X} of n values of X is normally distributed like $N(\mu_X, \sigma_X/\sqrt{n})$.

The Central Limit Theorem




Figure 1: The normal distribution about a mean μ with the units of standard deviations shown.

Example: An unknown distribution has a mean of $\mu = 90$ and a standard deviation of $\sigma = 15$. Samples of size $n = 25$ are drawn randomly from the population. Find the probability that the *sample mean* is between 87 and 93.

Interactive Questions

Interactive Questions: Central Limit Theorem

Suppose we take samples of size $n = 16$ from a large data set and compute the averages and standard deviations of the samples. Suppose we repeat the whole process, but change $n = 100$. Which of the following is true?

- A: The means of our samples will shift upwards by a factor of $100/16$.
- B: The means of our samples will shift downwards by a factor of $100/16$.
- C: The standard deviations of our samples will shift downwards by a factor of $\sqrt{100/16}$.
- D: The standard deviations of our samples will shift upwards by a factor of $\sqrt{100/16}$.

Interactive Questions: Central Limit Theorem

Suppose we take samples of size $n = 100$ from a large data set that has mean μ and standard deviation σ , and compute the averages and standard deviations of the *samples*. Which of the following is true?

- A: Each standard deviation of each sample we collect will be $\sigma/10$.
- B: The standard deviation of the means of our samples will be $\sigma/10$.
- C: Each mean of each sample we collect will be $\mu/10$.
- D: The mean of the means of our samples will be μ .

Confidence Intervals

Confidence Intervals

Suppose we are attempting to measure a population parameter, such as the mean, μ . This is distinct from the *sample mean*, \bar{x} , of a set of data $\{x_i\}$.




Figure 2: Let the gray curve refer to a *population* of data. Let the red curve refer to the *distribution of averages* we find if we repeatedly sample. The standard deviation of the red curve is smaller because of the Central Limit Theorem: $\sigma_{\bar{x}} = \sigma_{pop}/\sqrt{n}$.

Confidence Intervals




Figure 3: Suppose we decide to communicate the range in which the true mean μ falls, using our *point estimates* \bar{x} . How should we do that?

A good range: (Mean of \bar{x} distribution minus two standard deviations), to (mean of \bar{x} distribution plus two standard deviations). **This interval has a 95 percent chance of containing μ .**

Confidence Intervals

(Mean of $\bar{x} - 2s/\sqrt{n}$), to (Mean of $\bar{x} + 2s/\sqrt{n}$). **This interval has a 90 percent chance of containing μ .**

- The error bound for the population mean (EBM) is the “two standard deviations” in this example, in which we know the population σ .
- Recall the 68 - 95 - 99.7 rule, and how sample means are normally distributed regardless of the underlying distribution.
- We choose the confidence level (CL), and construct the confidence interval for the given CL.
- $\alpha = 1 - CL$, probability the confidence interval *does not contain* μ .

Professor example, next slide.

Confidence Intervals

Professor example.

PhET: Illustration of Confidence Intervals

PhET: Confidence Intervals

Go to this link: https://digitalfirst.bfwpub.com/stats_applet/stats_applet_4_ci.html

1. Select a confidence level (CL) of 95 percent using the slide bar at left.
2. Set your sample size to 100. The sample means will be distributed like $N(\mu, \sigma/\sqrt{n})$. This example assumes we know σ , the population standard deviation, but we do not know the mean μ .
3. Hit the green Sample 25 button to collect 25 samples of size 100, to get 25 means.
4. How many of these means fall within two standard deviations $\sigma/\sqrt{100}$ (CL: 95 percent) of the true mean μ ?
5. Repeatedly hit Sample 25, and notice what happens to the percentage at left of sample means that fall within $2 \sigma/\sqrt{100}$ of μ .

Interactive Questions

Interactive Questions

Suppose you wanted to measure the average age of inhabitants of a large neighborhood. The population mean μ is unknown, but we can be reasonably certain that our sample standard deviations, s_1, s_2, \dots are approximately σ , the standard deviation of the population. So we know that the $\sigma_{age} = 20$ years. Suppose we collect 100 samples of 16 people apiece. This means the standard error in our mean $20/\sqrt{16} = 5$ years.

1. Suppose the mean of all our data is 70 years. What is the EBM for two standard deviations in the mean?
2. How many of our samples should fall within 1 EBM of the true mean?

Solve via Breakout Rooms! Go.

Hypothesis Testing

Hypothesis Testing

The Null Hypothesis and Type I and II Errors

Let the **null hypothesis** H_0 be a claim, true or false, about a population of data, and H_a represent an alternative. If samples from the data confirm H_0 , then we accept H_0 as correct. There are two potential errors: Type I and type II. They are listed in Tab. 1. If H_0 and H_a are *complementary*, then rejecting H_0 confirms H_a .

	H_0 is True	H_0 is False
Do not reject H_0	Correct	Type II Error
Reject H_0	Type I Error	Correct

Table 1: Summary of Type I and Type II errors. The probability of a Type I error is α , and the probability of a Type II is β .

Hypothesis Testing

Steps for rejection or confirmation of the null hypothesis:

assume underlying distribution is normal via the Central Limit Theorem.

1. The probability of a Type I error is α , and the probability of a Type II is β . *Choose a value of α that is appropriate.*
2. Collect the sample of data. Calculate the mean (or other point estimate of a population parameter).
3. Determine the probability p that the mean occurred randomly, given the underlying distribution.
4. If $p < \alpha$, reject H_0 . If $p \geq \alpha$, accept H_0 .

Hypothesis Testing

Steps for rejection or confirmation of the null hypothesis:
assume underlying distribution is normal via the Central Limit Theorem.

Example procedure:

Hypothesis Testing




Figure 4: The square-wave signal spends equal time at 0.0 and 1.0, and the probability density function reflects that.

Hypothesis Testing




Figure 5: The triangle-wave signal spends equal time at all values *between* 0.0 and 1.0, and the probability density function reflects that.

Hypothesis Testing




Figure 6: The random noise *usually* spends time near 0.0, but rarely it fluctuates to larger values.

Example procedure for signals hidden in noise:

Hypothesis Testing

When a new drug is created, the pharmaceutical company must subject it to testing before receiving the necessary permission from the Food and Drug Administration (FDA) to market the drug. Suppose the null hypothesis is “the drug creates harmful side effects in 5 percent of patients.” What is the Type I Error?

- A: To claim the drug harms fewer than 5 percent of patients, and it harms fewer than 5 percent of patients.
- B: To claim the drug harms fewer than 5 percent of patients, and it harms more than 5 percent of patients.
- C: To claim the drug harms more than 5 percent of patients, and it harms fewer than 5 percent of patients.
- D: To claim the drug harms more than 5 percent of patients, and it harms more than 5 percent of patients.

Hypothesis Testing

When a new drug is created, the pharmaceutical company must subject it to testing before receiving the necessary permission from the Food and Drug Administration (FDA) to market the drug. Suppose the null hypothesis is “the drug creates harmful side effects in 5 percent of patients.” What is the Type II Error?

- A: To claim the drug harms fewer than 5 percent of patients, and it harms fewer than 5 percent of patients.
- B: To claim the drug harms fewer than 5 percent of patients, and it harms more than 5 percent of patients.
- C: To claim the drug harms more than 5 percent of patients, and it harms fewer than 5 percent of patients.
- D: To claim the drug harms more than 5 percent of patients, and it harms more than 5 percent of patients.

Hypothesis Testing

Suppose we are presented with the data in Fig. 7. H_0 : there will be no outliers more than three standard deviations from the mean in 256 days. Do you accept H_0 ? Why or why not? What is your chosen α ?




Figure 7: The change in the account, in USD, versus time, in days.

Conclusion

Unit 2 and 3

1. Central Limit Theorem: 7.1
2. Confidence Intervals and Hypothesis Testing
 - Confidence intervals and data interpretation: 8.1 - 8.4
 - Rejecting the null hypothesis, types of error, underlying distributions: 9.1 - 9.3